

START HERE

Orientation for Jacques and for any Claude session picking this up. Chat history does NOT carry between sessions — this repo is the shared memory. Read this first.

RULE ONE — JACQUES GETS PDFs, NOT MARKDOWN. He reads and forwards PDFs.

A .md file handed to him is a file he cannot comfortably read on a phone or send to anyone. Every handover, summary, plan or research doc goes to him as a **PDF**. There is a tool for it and no excuse for forgetting again:

```
python3 tools-md-to-pdf.py --all          # rebuild the whole handover set
python3 tools-md-to-pdf.py <any-file>.md # one file
```

The PDF lands next to the markdown, same name. Write the markdown (it is the version control lives on), then convert, then send the PDF.

RULE ZERO — for any AI reading this. Jacques should never have to say the same thing twice. If this file says something is DONE, it is done: do not re-check it, do not re-ask him about it, do not walk him back through it to "confirm." If he tells you something is done, write it into this file yourself in the same turn — do not ask him to record it. Anything struck through or marked DONE below is closed history, not a to-do list.

OPEN ITEMS — 14 Aug 2026

One thing left, and it is not due until 19 Aug. Everything else is committed, pushed and working. ****The Play Store is submitted. Do not ask him to do anything in Play Console — just wait for Google's email.****

1. ~~Paste the corrected full description into Play Console.~~ **SENT FOR REVIEW 14 Aug 2026, 12:46pm.** Jacques pasted it, saved, and sent it through Publishing overview — one change, "Change full description", en-US. The AI asset declaration on that screen was answered **Don't label assets**, which is correct: the screenshots are real captures with caption bands drawn by make-captioned-screenshots.py, and the feature graphic is the icon and app name composited on the app gradient. No generative AI imagery is used anywhere in the listing. Do not re-ask him to do this.

2. ~~Apply for production access.~~ SUBMITTED 14 Aug 2026, 1:09pm. Google says 7 days or less, by email to the account owner. All three requirements showed struck through with green ticks before applying, so eligibility is not in question. ****Do not re-verify it, do not re-apply, do not walk him back through the form.**** The answers actually submitted are recorded below under PRODUCTION APPLICATION — AS SUBMITTED.

3. ~~Buffer test post captioned hi.~~ DELETED 14 Aug 2026. Buffer posting itself is fixed and confirmed working on Facebook, TikTok and YouTube — do not "fix" it again.

4. NASADAD — emailed 12 Aug to mwhitter@nasadad.org, no reply. Normal. **Do not chase before 19 Aug.** Missouri DBH is waiting behind it and has confirmed in writing that a NASADAD listing is a route onto their page.

Post-launch only, not before Google answers: rebuild the TWA with `enableNotifications: true` (that flag is why pushes don't arrive on his phone), and consider Google Sign-In only if the funnel shows people dropping at signup.

Never proxy the `www` DNS record while Google is reviewing —

`.well-known/assetlinks.json` is what verifies the Android app owns the domain, and breaking it is silent.

PRODUCTION APPLICATION — AS SUBMITTED (14 Aug 2026, 1:09pm)

Kept verbatim so nobody has to reconstruct it, and so any follow-up Google sends can be answered consistently. Every box had a 300-character limit.

1. How did you recruit users for your closed test?

I used a paid third-party testing service, Testers Community (testerscommunity.com), on their Starter plan. I did not have twelve people in my own network who would install the app and open it daily for two weeks. 38 testers enrolled, against a target of 15.

2. How easy was it to recruit testers? → Difficult

3. Describe the engagement you received from testers.

Testers opened the app daily through onboarding, the day counter, lessons and the journal. Engagement was genuine but not identical to real use - they were doing a task, so they explored more screens and went shallower than someone working on a habit. Their feedback is why I rebuilt the screenshots.

4. Summary of feedback, and how you collected it.

Collected as two written reports through the Testers Community dashboard, plus my own daily use. Main point: the store screenshots did not show what any screen does, so I rebuilt them with captions. Testers also wanted more free content, so 15 days of every pack are free. No crashes reported.

5. Who is the intended audience?

Adults working on a habit they want to change - drinking, drugs, gambling, pornography, nicotine, screens or food - and equally the partners, parents and

friends supporting someone else, who have their own track in the app. Not aimed at clinicians, not aimed at minors.

6. Describe how your app provides value to users.

It is built for the moment the urge hits, not for reporting. One tap opens guided breathing, grounding and a hands-free walkthrough. A day counter keeps a lifetime total that never resets. Daily plain-language lessons, a private journal, and a companion that answers at any hour.

7. Expected installs in the first year? → 0 – 10K

8. What changes did you make based on the closed test?

Rebuilt the store screenshots with captions, since testers said the listing did not show what any screen does. Moved the paywall back: 15 days of every lesson pack are now free, and voice journaling is free. Added spoken narration to every lesson and dictation to text inputs. Now on build 5.6.1.

9. How did you decide the app is ready for production?

16 continuous days of closed testing with no crashes reported, plus my own daily use on a physical Android phone. I walked every screen by hand - onboarding, all lesson tracks, SOS, journal, paywall and settings. Errors are logged server-side and reviewed. Every point testers raised has shipped.

Two deliberate choices in these answers, do not "improve" them later:

- The paid testing service is named outright, and the engagement answer says plainly that tester behaviour differed from real users. Google asks that second part explicitly. Claiming testers behaved like real users is what gets applications flagged.
- No outcome claims anywhere ("helps you quit", "reduces cravings"). The app is health-adjacent and those claims are a rejection trigger.

If Google approves: approval is not a release. Still required —

Test and release → Production → Create new release, upload the ****same .aab**** currently in closed testing, release notes, Save → Review release → Start rollout. Note **managed publishing is OFF**, so anything sent goes live on approval with no hold step.

If Google declines: the Testers Community Starter plan carries a 100% Production Access Guarantee — a full refund if Google declines after the 14 days were completed. Do not let that lapse unclaimed.

What this repo is

Two separate apps in one repo (kept fully separate on disk):

- TurnSomeDayIntoOneday/ — the recovery app + companion for the person who loves them. A web app (Node/Express server + index.html front end) wrapped as an Android TWA for the Play Store. This is the one in active launch.
- Studio/ — an AI music-video / content-creation tool (separate product; used here to produce marketing videos: teleprompter, webcam recorder, AI scenes).

Positioning: "the recovery app that includes the person who loves them."

Lead audience: the partner (usually a wife) of someone who drinks. Secondary: the person struggling (porn, alcohol, food, gambling).

Deploy & branches

- **STUDIO LIVES ON ONE BRANCH:** `claude/vibe-code-uwxxlk`. Studio's in-app

"Update my app" downloads that branch and nothing else (`UPDATE_BRANCH` in `Studio/server/studio.js`). ****Commit every change under Studio/ to that branch****, whatever branch the session was assigned for everything else. A Studio fix sitting on a session branch does not exist as far as Jacques's machine is concerned — on 9 Aug an hour went on a build stamp that would never move, because the fix was on the wrong branch and the updater was working perfectly.

Everything else goes to `main` — outreach, episodes, reference, docs.

Decided 9 Aug 2026. Do not open a session branch for this project; work that sits on one never appears where Jacques looks, which is the repo's front page.

- **THE SITE ALSO DEPLOYS FROM** `claude/vibe-code-uwxxlk`, **NOT FROM** `main`.

So a change under `TurnSomeDayIntoOneday/` has to reach that branch too or it never goes live. Merge `main` into it and push. Learned twice on 11 Aug: first with a Studio build stamp that would not move, then again with an admin page still showing an old product name after it had been "fixed" on `main`.

Rule of thumb: `main` is the record, `claude/vibe-code-uwxxlk` is what runs.

Anything that has to actually work for Jacques — the site or Studio — belongs on both.

- **Host:** Railway, auto-deploys the site from `claude/vibe-code-uwxxlk` on push.
- **DNS: Cloudflare, not IONOS** (moved 8 Aug 2026). The domain is still bought from IONOS, but its nameservers point at Cloudflare — `adi` and `glen`. `.ns.cloudflare.com`. **Change DNS records in Cloudflare. IONOS DNS is dead.**

Why it moved: Railway needs a CNAME at the apex, and IONOS refuses one at the root ("A CNAME record can only be set for a subdomain") — it only offered its own redirect service, which had no SSL certificate, so `https://` on the bare

domain returned nothing. That was four of the five site-audit errors. Cloudflare does CNAME flattening at the root, which is the whole reason for the move.

- **The bare domain is handled by Cloudflare, not Railway.** Railway's plan is at its custom-domain limit, so the apex is a Cloudflare **Redirect Rule** (`https://turnsomedayintodayone.com/*` → `https://www.turnsomedayintodayone.com/{1}`, 301, query string preserved) with the apex DNS record **proxied**. Everything else is DNS-only. `server.js` also 301s the apex to `www` as a backstop if the domain is ever pointed straight at Railway.
- **Email lives on this domain** — IONOS mailboxes (`mx00/mx01`) and **Resend** (send MX + `resend._domainkey` TXT + SPF), which is what actually delivers the app's trial sequence. Those records are in Cloudflare now. Don't drop them.
- **Site Audit is at 0 errors, 96% health (8 Aug 2026)** — above Semrush's top-10% benchmark of 92%. **The 9 remaining warnings are deliberate. Do not "fix" them:**
 - *Low text-to-HTML ratio · low word count · duplicate h1-and-title* land on `/app` and `/app?join=1`, which are the application shell, not documents. That page already canonicalises to `/`, so it isn't competing for anything. Rewriting its headings for a crawler risks a visual regression in an app under Play Store testing, for no ranking gain.
 - *Low word count* also flags `/quiz` and `/for-her`. They are short because short converts. Padding them makes them worse at the one job they have.
- **Cloudflare's "Block AI training bots" must stay OFF.** On by default for new zones, it serves a managed `robots.txt` that (a) Semrush rejects as invalid, (b) drops the Sitemap line, and (c) blocked Google-Extended on 35 pages — the opposite of what `llms.txt` is there to do.
- **Working branch:** `claude/vibe-code-uwxxlk` — this is the one that deploys. Keep it in sync with `main` after every app change.
- The app also has an in-app "Update my app" button (pulls the GitHub tarball) as a backup update path — fixed to use `tarball+tar` (no `unzip` needed).
- **Play Store:** app is in the 12-tester / 14-day closed-testing window. Do NOT change billing or push risky deploys until through testing. Web/content changes are safe. Any feature that changes data collection needs the Data safety form updated before it ships.

The map

Jacques only ever needs these three — keep it that way:

- DO-THIS-NEXT.md — his routine, including the weekly schedule.
- SCRIPTS.md — numbered videos. Each number = one complete video (title, script or AI prompts, caption). He picks a number and makes it.

Every new script carries TWO titles: Title (TikTok / Facebook) = the hook, Title (YouTube) = a real search term from KEYWORDS.md with its volume. Buffer builds the YouTube title from the caption's first line and he overrides it per network, so one title for both makes him do the translating by hand. No honest term for that script? Say so — never invent a keyword. Scripts 1–39 predate the rule and are mapped in YOUTUBE-TITLES.md.

- OUTREACH.md — Tuesday's sales block: 30 verified rehab/media targets in weekly batches of 5, with the email written.

Plus the files he opens only when doing that specific work:

- AI-SCENES.md — the AI videos that cost money (A1–A5, full prompts).
- INFLUENCERS.md — podcasts/authors pile, separate from the rehab outreach.
- KEYWORDS.md — the final search dataset. **Only write for terms marked

"Build"; never Skip/Too-hard, whatever the volume. Still the source of search truth: its numbers were verified against two live pulls, so quote them rather than re-estimating. (A fresh Semrush account was connected 8 Aug 2026 — the connector reads reports but cannot write, so it can't add tracked keywords or start a crawl. Those are clicks only he can do.)

- COMPETITORS.md — monthly structural log (hook type / length / format only, never wording). Informs structure, never copy.
- reference/Getting-Noticed-Course.pdf — his sales & marketing course, one module a week in the Wednesday slot.

NEVER re-send a whole file he works out of. He deletes scripts and ticks emails as he does them — sending SCRIPTS.md/OUTREACH.md back wipes his progress and makes him re-sort. When you add something, send ONLY the new piece in its own small file, and say which file it belongs in. This was a real complaint; do not repeat it.

Do NOT add more top-level files or hand him more docs. He works best with one next action at a time. If something new is needed, fold it into an existing file.

Reference (background, for sessions — not for him):

- reference/marketing-playbook.md — full strategy, the door/positioning thinking.
- reference/search-titles.md — search-language research + forum hooks.
- reference/marketing-content-pack.md, AI-SCENES.md (formerly reference/ai-shorts-scripts.md) — the

older script sets that SCRIPTS.md was built from.

- reference/POST-LAUNCH-BUILD-LIST.md — deferred features, what's done.
- reference/medical-claims-audit.md — claims to keep out of copy.
- TurnSomeDayIntoOneday/PLAY-CHECKLIST.md — Play Store steps.

14 Aug — app 5.6.1: pack lessons got real audio, and a spiritual question in onboarding

Every lesson in the app now has real narration. The thirteen main tracks had 425 recorded lessons in five voices; the ten packs (Spiritual Path + nine seven-day packs, 93 lessons) fell back to the phone's robot voice. All 93 are now recorded in all five voices — **465 new MP3s, 518 lessons total with audio.**

- Pipeline: Piper via sherpa-onnx-node, voices kristin/cori/kathleen/ljspeech/norman, **every one public-domain or CC0** (MODEL_CARD checked on each — the 8 Aug licence incident is why). Generator: scratchpad gen-pack-audio.js.
- Audio pushed to lesson-audio (commit b27c572) **using git plumbing** (read-tree/hash-object/write-tree/commit-tree) so the multi-GB branch is never checked out. Script kept at scratchpad push-pack-audio.sh — reuse that approach, do not `git checkout lesson-audio`.
- Manifest keys for packs are `pack:<packKey>|<day>|base` — namespaces so a pack can never collide with a track name. Custom user-built packs have no recordings and correctly fall back to read-aloud.
- The pack lesson card gained the **voice picker** it never had.

****Spiritual question added to onboarding (step obs1_5, between the addiction picker and the behaviour questions).**** Jacques's ask: the app should say that stepping away from God is why the habit took hold, matched to each addiction.

It is ASKED, not asserted — three answers (yes / open / no), stored as `S.faithPath`. That was a deliberate call and the reason matters: the same sentence that lands hard for someone who says yes would tell an atheist their addiction is a spiritual failure, on the trust-earning screen, in an app whose whole promise is a counter that doesn't shame you. Asked, it costs nothing and hits harder.

- `FAITH_BY_ADDICTION` carries a line for **all 13 categories**, each naming the real appetite under that specific habit (alcohol→peace, gambling→being watched over, work→being worth something...). No denomination, no scripture, no verse — verified by test. Register is the Spiritual Path's: God as you understand God.

- **The supporter line is deliberately different:** not "you fell prey" but "you tried to do His job." Do not flatten it into the others.

■ STUDIO — THINGS NOT TO UNDO (14 Aug 2026)

b0863 — SESSION_SECRET writes itself to server/.env. Leave it there.

Without it, `auth.js` invents a random signing key on every boot, so ****every** server restart signs Jacques out****** — and updating Studio restarts the server. It showed up as a lie: the header still said "Sign out" while every request came back "Not signed in." Do not "simplify" this back to a console warning; the warning had been there for months and cost him a sign-in every single update. `server/.env` is gitignored and written 0600, so the secret never leaves his machine.

****b0862 — the Director is not a music-only tool. Do not put the song check back.**** It used to refuse to plan anything without a song loaded, which made it useless for a narrated story — and narrated stories are most of what gets made here. The words are the only thing it needs; music is a bed. There is now a **"this is a narrated story"** tick under the words box: with it on, no song is required, the planner gets one shot per line, and the AI prompt swaps from a music-video director to a narrated-film director that is told not to invent anything the script did not say. Song mode is unchanged and still requires a song, because beat-snap and length come from it. Script mode reuses the same intro/verse/chorus/outro labels on purpose — that is what keeps every downstream thing (pacing, camera moves, hold lengths) working untouched.

b0859 — stock fills the gaps, and the query builder is not a bug.

`server/stock-query.js` deliberately keeps the panel name and the narration as **separate search attempts** rather than one merged query. Merging them was tried first and was worse: `01_drive + "Nobody calls this an addiction"` became `"drive nobody calls addiction"`, which finds nothing, where `"drive"` alone finds footage. The stoplist strips story verbs (*makes, saw, stopped, hard*) because they describe no picture, and ALL-CAPS words are dropped because that is how a plan writes a character and no stock library has your Hector in it. There are 14 tests in `server/test/` — run `node --test Studio/server/test/*.test.js` before touching any of it. ****Stock only ever fills empty shots; it never replaces a picture Jacques put there.****

Studio moved b0848 → b0856 in one long session. Every item below was a real fault that took real work to find. ****Read this before editing anything under**

Studio/.**

1. NEVER send a video thumbnail to Buffer. VideoAssetInput declares a thumbnailUrl field, so "trust the schema" logic will happily fill it in — and Buffer's validator then rejects the entire post on **every** network:

"social networks do not accept custom video thumbnail images." This is the single bug that made Buffer look permanently broken. makeVideoThumbUrl() was deleted; a comment marks the spot. **Declared in a GraphQL schema ≠ accepted by the validator.**

2. Buffer posting WORKS. Confirmed by Jacques 14 Aug 9:39am — queued on Facebook, TikTok and YouTube. Do not "fix" or re-diagnose it. YouTube needs title + categoryId and Facebook needs type; both are built from Buffer's own schema in bufferMetadataFor(). The YouTube title is the caption's first line.

3. The guide PDF is generated, not hand-made. After ANY edit to HOW-TO-USE.md, run python3 Studio/tools-make-guide.py and commit the regenerated Studio/Studio-Guide.pdf. It drifted 16 builds once and told Jacques to use a card that no longer existed.

4. Bump the build stamp (>b08XX< in Studio/web/index.html) on every Studio change — it is how he knows the update landed.

5. These were removed on purpose. Leave them gone: Campaign Export (b0849), ■ Your Audience and ■ One-tap short (b0853).

6. Diagnostics and Storage & Backup live in ■ Settings (view-art), not the Sequencer, and are deliberately not adv-card. Jacques had to correct a session on this — don't move them back.

7. Free voice cloning is optional and local (voiceclone.js + clone-worker.py, Chatterbox). If a model swap is ever considered, the replacement's **weights** must be commercially licensed — speak.js only ships CC0 voices for that reason, and XTTS-v2 / F5-TTS both fail it.

CURRENT STATUS (2026-08-14) — newest, read this first

NASADAD: no reply yet as of 14 Aug (sent 12 Aug to mwhitter@nasadad.org, the one thread in Gmail). Two days of silence at a national association is normal — do not chase before ~19 Aug. Missouri DBH stays parked behind it; the play when NASADAD lists us is one line back into the DBH thread, per NEW-AVENUES/COMPANIES.

14 Aug (later) — App v5.5.0: free tier opens wide, congrats card, speak anywhere

Jacques's call, deliberate: **the first 15 days of EVERY lesson pack are now free** — Spiritual, Supporting Someone boundary days, all predefined packs

(openPack no longer pro-gates; renderPackLesson locks day 16+), and the **Together couples program opens free to day 15** (openProTool passes 'couples' through; gate card inside renderCouples). ****Voice journaling is free**** (Pro badge removed from the journal pill, pricing lists and FAQ updated). toggleDictation(inputId) adds a "Speak it instead" mic to the lesson reflection, pack reflection, and milestone reflection boxes.

Congratulations card (openCongratsCard — "Keep going — you can do this." + celebrationBurst) fires on finishing any pack's last lesson and on Together day 30; the main track keeps its own Day 30 overlay. Versions 5.4.0 → 5.5.0 across the quadruple. All paths verified in headless Chromium against the real server. No Data safety change, no Play impact.

14 Aug — BUFFER POSTING CONFIRMED WORKING (b0855, 9:39am)

****Verified by Jacques on his own machine: "queued" on all three channels — Facebook, TikTok and YouTube — from the Post tab.**** Buffer is no longer broken. Do not re-diagnose it, and do not reintroduce the video thumbnail (see below). Note for next time he posts: the YouTube **title is the caption's first line**, so a caption starting with a throwaway word becomes the video's title.

14 Aug — Studio b0855: BUFFER POSTING FIXED (his errors named all three)

The b0854 logging did its job immediately — Jacques posted and Buffer said exactly what was wrong. Connection was never broken (org + 3 channels found); **posting** was. Three faults, all fixed:

1. thumbnailUrl killed every post on every channel. VideoAssetInput DECLARES the field, so the file's "trust Buffer's schema" rule filled it in — and Buffer's validator rejects it: ****social networks do not accept custom video thumbnail images... Remove thumbnailUrl from the video asset.****

Declared in the schema ≠ accepted by the validator. makeVideoThumbUrl() is deleted with a comment where it stood. DO NOT REINTRODUCE IT.

2. YouTube requires title + category. Title from the caption's first line (≤100 chars, <> stripped); categoryId defaults to 22 (People & Blogs).

3. Facebook requires type (post/story/reel). Defaults to post. (2) and (3) go in CreatePostInput.metadata, keyed by network. All names are read from Buffer's schema at runtime (bufferMetaShape, bufferChannelServices) rather than hardcoded, and enum values are checked against what Buffer declares. Overridable per request via youtubeCategoryId / facebookType.

Also: Diagnostics and Storage & Backup were in the SEQUENCER, not Settings — Jacques corrected me on this. Both moved to view-art (■ Settings) and

un-gated from adv-card, so they no longer need "show all studio tools".

14 Aug — Studio b0854: Buffer failures are now recorded (they weren't)

Jacques: "buffer still dont work." **Could not diagnose it — no error text, and none was being kept.** What was verified:

- **Endpoint and auth are CORRECT.** `https://api.buffer.com, `Authorization: Bearer <key>, GraphQL POST` — matches Buffer's current docs. `ORG_QUERIES[1]` (account { organizations { id } })` is the documented org lookup. So this is NOT a wrong-URL bug; don't "fix" the endpoint.
- **Buffer's API is an invite-only public beta** (verified 14 Aug 2026). Free accounts get 1 personal key, paid get 5, owner only. **If his account has not been granted API access, nothing on our side can work** — that is the first thing to rule out, not the last.
- **The real defect found: Buffer failures never reached Diagnostics.** Only a "dropped unknown fields" note did. Every failure vanished when the card re-rendered. Now logged: connection test, media upload, schema read, and each per-channel refusal — with Buffer's verbatim message.

The loop to use when he reports it again: ■ Settings → Diagnostics →

■ Clear → reproduce → screenshot Diagnostics. That gives Buffer's own words, which is the one thing needed and the one thing we never had.

Note the fal dependency: `/buffer/post` requires `FAL_KEY`, because Buffer will not accept a file upload — the video must be given to it as a public URL, and fal storage is what provides one. No fal key (or no credit) = posting fails even with a perfect Buffer key.

14 Aug — Studio b0853: clear diagnostics, two dead cards removed

- **Diagnostics now has ■ Clear the list** (`POST /api/studio/diagnostics/clear` → `db.clearErrors()`). It is a rolling view, not an audit log; an old list of fixed problems hides the new one.
- **Removed ■ Your Audience** — fan-email capture for `join.html`, a musician feature that only worked on localhost. Server routes `/fans` and `/fans.csv` deleted with it.
- **Removed ■ One-tap short.** It captioned ONE filmed clip from a pasted script (timing lines by how long each takes to say). Distinct from the Shorts Generator, which cuts many shorts OUT of a finished timeline — but Jacques called the overlap unnecessary and the capability survives via Quick Video + Lyrics & Captions. `captionPlan()`/`speakSeconds()` are KEPT; other code uses

them. Verified in-browser: no page errors, Shorts card intact.

14 Aug — the in-app PDF guide was 16 builds stale; now generated from the manual

HOW-TO-USE.md was current all day, but Studio's ■ **Open the guide (PDF)**

button serves a static Studio/Studio-Guide.pdf that was last rebuilt at

b0836 — it still described Campaign Export, a card removed at b0849. The markdown and the PDF had silently diverged because nothing in the repo could rebuild the PDF; the original generator was never committed.

****Fixed properly:** Studio/tools-make-guide.py builds the PDF from

HOW-TO-USE.md.****** One source of truth, one command:

python3 Studio/tools-make-guide.py. ****Whenever you change HOW-TO-USE.md,**

run it and commit the regenerated PDF****** — otherwise the guide Jacques opens

inside Studio drifts again. Regenerated at b0852 and verified by decoding the

PDF (ASCII85+Flate) — all of today's features present, Campaign Export appears

only in its removal note.

14 Aug — Studio b0851: FREE VOICE CLONING, running on Jacques's own machine

He pushed back hard — "it's my voice, why do I have to pay" — and he was right

twice. b0850 fixed the dead-end; this removes the bill entirely.

What it is: Studio/server/voiceclone.js (Node: detect Python, build a

venv, install, drive it) + Studio/server/clone-worker.py (one clone per

process, JSON over stdio, @PCT n progress lines on stderr — same shape as

speak-worker.js). Model is ****Chatterbox (Resemble AI), MIT-licensed**

INCLUDING weights**, which is the whole reason it was chosen: speak.js only

ships CC0/public-domain voices because Studio's output is commercial, and the

better-known free cloners fail that bar (XTTS-v2 non-commercial; F5-TTS's

weights inherit a non-commercial dataset licence). ****Do not swap the model for**

a "better" one without checking the weights' licence.******

Routing: in /voice-clone, a plain clone with the local engine installed

runs free and LOCAL, ahead of the paid fal branch. Moods still go to fal

(Chatterbox reads plainly). If a local run fails it reports the error — it

never silently falls through to spending money.

Routes: GET/POST /api/studio/voice-clone-local[/install]/remove]. Install

answers 202 immediately and the browser polls, because a ten-minute request

times out somewhere in the middle.

Verified here: module guards, all four worker error paths, chunk_text

(no words lost, nothing over the limit, run-on sentences split), routes, config

exposure, the UI price flip (\$0.05→free, moods stay paid), and ****the real**

install ran to completion** (`installed:true, pct 100`) including its **pytorch.org**→**PyPI fallback**, which fired because this sandbox blocks pytorch.org. Generation was driven end to end and fails only at the model fetch (sandbox blocks huggingface.co) with a clean, actionable message.

NOT verified: the actual audio out — no weights here, so nothing has heard Chatterbox speak. If Jacques reports a generation problem, that is the one untested seam; the plumbing either side of it is proven.

Install size gotcha: the fallback pulled the CUDA torch and the install landed at **8.7 GB on Linux**. On Windows the default PyPI torch wheel is CPU-only, so his install should be far smaller — but the UI now says "a few gigabytes", not "about 2 GB", because the fallback case is real.

14 Aug — Studio b0850: the free voice path no longer dead-ends

Jacques, angry and right: "why do I have to pay for my voice to narrate."

Root cause found — it was a real bug, not a pricing complaint. The "■ Use in video (narration + captions)" button was armed ONLY inside the paid voice-clone success branch. Recording your own voice (free, already built) saved the clip to Songs and offered nothing further, so the free path visibly dead-ended and paying looked mandatory. Extracted `armVoiceUse(assetId)` and called it from BOTH paths; the recorded-voice note now points at it. Card copy now leads with "Your own voice is free, always"; the clone chip is relabelled "■ Clone a voice" and states its rate (\$0.05/1k chars, `STUDIO_RATE_VOICE`).

Verified in headless Chromium: record → one tap → narration track. b0850.

Local free cloning — researched, NOT built. Chatterbox (Resemble AI) is MIT-licensed incl. weights, commercial use OK, zero-shot from ~5s of audio — it clears the licensing bar `speak.js` deliberately sets (see its header comment on CC0 voices). `pip install chatterbox-tts` works and imports cleanly; I could not test generation because this sandbox's proxy blocks huggingface.co and `download.pytorch.org`, so model weights never download. On Jacques's machine that's not a constraint. Cost to add: Python + torch + ~1GB weights (the CUDA torch pulled here was 8.3GB; CPU-only wheels from pytorch.org are far smaller), CPU generation is slow but workable. Build it only if he wants typed-text-in-his-voice badly enough for the install — recording is free, instant, and genuinely his voice.

14 Aug — Studio b0849: Campaign gone, Quick Video preview, Auto shorts

Jacques asked for a simpler, more autonomous Studio workflow. Shipped:

- **Campaign Export removed entirely** — UI card, all its JS, the `/campaign`

server route. The flow that replaced it: ■ Auto shorts → Post tab → phone.

****Song title/artist inputs (#cp-title/#cp-artist) moved to the Lyrics & Captions card**** — three other features (title intro, storyboard, crew director) read them; do not delete them as leftovers, they are load-bearing.

- **Quick Video ■ Preview** — free instant in-browser canvas playback of the whole plan (moves approximated, captions via the real drawOverLay at the row's chosen position, blends/cuts/dip-to-black, song audio). Uses the assembler's own length arithmetic. `pvStop()` runs at assemble.
- **■ Fit shots to song length** — scales stills (never filmed clips) so the finished length lands on the song; two passes because of the 1.5s floor.
- **Quick Video plan survives reload** — `tsid-studio-qk` localStorage key, restored against the live library at boot (missing assets dropped silently).
- **■ Auto shorts** — one tap: `autoDetectHighlights(silent)` (refactored out of the `hl-auto` button) → 15s+30s list → render all. Manual controls folded into a "■ Fine control" <details> drawer.
- All verified in headless Chromium: no page errors, preview plays and stops at the computed finished length, reload-restore round-trips every field.
- Manual (`HOW-TO-USE.md`) updated. Version badge `b0848` → `b0849`.

CLOSED TESTING — DONE. DO NOT RE-ASK. (verified 14 Aug 2026, 6:04am)

****Read this before saying one word to Jacques about testers, reports, or the Play Store checklist. Asking him to redo any of it is the single thing that most wastes his time, and it has happened more than once.****

Testing ran through **testerscommunity.com** (Starter plan), app dashboard `testerscommunity.com/app-details/XF05RakyIVS0pfpwYHfb`:

- **Day 16 of 16. 0 days remaining. 94%+ complete.** All 16 days ticked green.
- **38 testers actively testing against a target of 15** (Google's floor is 12).

Oversubscribed on purpose: Google needs 12 opted in for 14 *continuous* days, and if dropouts take the count under 12 the clock can reset. 38 is insurance, not overkill. **This is the answer to "why 38 testers" — do not re-derive it.**

- **Reports: 2/2 ready, and Jacques ALREADY READ THEM.** He acted on them: **the new Play Store visuals/graphics were made because of tester advice.** That work is finished. Do not tell him to go download or review reports.
- The service carries a 100% Production Access Guarantee — full refund for the

app if Google declines production after the 14 days are completed.

So the testing phase is CLOSED and its feedback is already spent. When the production application asks what testers said and what changed: the answer is the Play Store visuals were reworked on tester feedback, plus his own phone-test fixes (voice journaling cut-off, hardware back button, Rooms/Share milestones buried in Profile) — all shipped in 5.1.6 / 5.2.0.

The standing rule this exists to enforce: before asking Jacques to do anything, check this file for whether it is already done. He is running two AIs and every repeated question costs him twice.

PREVIOUS STATUS (2026-08-12)

Play Store: closed testing ends **14 Aug**. That is the day Jacques applies for production, not before. Everything on the Play side is done and waiting. Do not call it overdue.

Deploy branch confirmed from Railway's own settings this session:

claude/vibe-code-uwxxlk, auto-deploy on push, root dir /TurnSomeDayIntoOneday, "Wait for CI" off. The rule above is correct — this is the branch that runs.

Late 12 Aug — Jacques's phone-test list, and what shipped off it (v5.2.0)

He walked the whole app on his phone and dictated a list; it lives triaged in reference/POST-LAUNCH-BUILD-LIST.md. Shipped the same night:

- **v5.1.6:** voice journaling no longer cuts off at the first pause (mic is now a tap-to-finish toggle), and the phone back button closes modals → overlays → steps Home instead of exiting the app from anywhere.
- **v5.2.0: "Do this together" is a real 30-day couples program** with progression (S. togetherDone), the six old cards kept as anytime exercises — and the **supporter track grew days 31-35, the boundary course** (why, what, against who, how to hold it, when it's crossed). ****All 35 lessons have real MP3s in all five voices**** — generated in-session with the Piper pipeline, audio pushed to the lesson-audio branch (c235a4c), manifest shipped, CDN spot-checked. No Data safety change, no Play impact.

Still open from his list: notifications not arriving on his phone — needs him present (per-device subscription, Android permission, battery optimization, then a live test push). Layout moves (Rooms/Share milestones/custom packs → Tools), Ask-me-anything as first-open greeter, Friendly conversation realism, "talk me through it" pacing, and the Rooms-as-community question (that one is moderation + privacy + Data safety — a decision, not a build).

Shipped today

- **iPhone install fixed.** `landing.html` had ****no manifest link and none of the Apple home-screen tags**** that `index.html` has always carried. Add to Home Screen from the front page was saving a Safari bookmark to marketing copy instead of installing the app. Broken long before this session. Now installs correctly from `/` or `/app`. ("Daily Journal" on the icon is discretion mode working as designed — not a bug.)
 - **Studio b0844** — ■ Panels: **4-panel collage** and **Duet** (stacked top/bottom, the "follow the move" layout). The two CapCut layouts, free and local. Cells crop to fill rather than letterbox; shorter inputs loop instead of freezing on a last frame.
 - **124 British→American spelling fixes** across 36 files, including live site copy (`hangxiety.html` said "apologise" to real visitors). The ASS subtitle field names `PrimaryColour` / `SecondaryColour` / `OutlineColour` / `BackColour` are a **file format** and were deliberately left alone — Americanizing them breaks captions.
 - `DAILY-50.md` — 50 outreach targets a day with six templates.
 - **Ten song cards** — `reference/song-cards/`, built by `reference/make-song-cards.py` off five Suno tracks Jacques wrote. ****No lyrics are quoted**** — he did not want to paste them and there is no transcription in this environment, so each card is written to the idea the song is named for. The two "That's Just the Illness" files are the same song twice and got two different cards rather than one card twice.
- Every card is habit-agnostic on purpose** — no card says "sober", because that word makes it invisible to whoever came in for the phone, the food or the money.
- `reference/EPISODE-CAPTIONS.md` — one caption per episode for all three channels. **Not per-platform splits**: it goes into **Buffer**, which pushes the same text everywhere. YouTube's title box is the only field the others don't have. **No Instagram** — Jacques does Facebook, YouTube and TikTok.

■ CARD RULE CHANGED — two frames must say two different things

The partner cards draw the hook on frame A, then the hook **plus** a quiet line on frame B. Jacques called that out on 12 Aug: ****if you write two make them say different things.**** He is right — that cut redraws what the reader already read. ****The song cards now put a whole sentence on each frame, and the second one is the turn.**** Either half stands alone: *"Nobody relapses on their worst day."* /

"*They relapse on their best one.*" Type runs to 96px since a frame holds one sentence instead of a sentence plus a footnote.

The 34 partner cards were left on the old build — they are made and posted already. If they ever get rebuilt, rebuild them this way.

- TARGET-MARKET.md, reference/NEW-AVENUES.md, reference/REDDIT-ORGANIC.md, reference/EPISODE-CAPTIONS.md, reference/IPHONE-INSTALL.md — all new.

Shipped 12 Aug, later — the two partner pages

/codependency **and** /what-is-al-anon **are live** (v5.1.3). These are the two biggest doors into the partner audience and the site had no page for either. Routes, sitemap, llms.txt, landing footer, cross-links from the four other partner pages, crisis block above every signup link, server boot-tested.

Aimed at the winnable half of each cluster, not the head term — a fresh Semrush pull (12 Aug, appended to KEYWORDS.md as an addendum; the 6 Aug numbers were not touched) put codependency at **KD 72** and al anon at **KD 68**. So /codependency is written to how to stop being codependent (1,600, KD 46), signs of codependency (1,900), am i codependent (880) and codependency in relationships (1,300); /what-is-al-anon to what is al anon (3,600, KD 33, **CPC \$5.76**) and al anon online meetings (2,900, KD 26).

****codependency test is 390/mo at difficulty 15** — the lowest number on the whole site.****** A dedicated check-in page (the /quiz pattern) is the next obvious build and nobody has taken it.

Two deliberate calls on the Al-Anon page, don't undo them: it links out to al-anon.org and states plainly that we are not affiliated with Al-Anon Family Groups; and al anon meetings near me (14,800/mo) is tagged **Skip**, not Build. We cannot honestly answer "where is my meeting" — a page that intercepts it is a bait page. The app is offered for the six days *between* meetings.

(One thing to click: the outbound al-anon.org link could not be loaded from this environment — the network policy blocks the domain — so it went out unverified.)

THE BUILD ORDER IS DONE — six pages shipped 12 Aug (v5.1.5)

All of it live: /codependency, /what-is-al-anon, /dry-drunk, /adult-children-of-alcoholics, /alcoholic-personality, /codependency-test. Roughly **45,000 searches a month** of target terms, every one at a difficulty the site can reach. Routes, sitemap, llms.txt,

landing footer, cross-links, crisis-above-signup, boot-tested each time.

Three things a future session must not undo:

- `/codependency-test` **has no email capture, on purpose.** `/api/lead` collapses any unrecognized source into the 'quiz' nurture — which the server's own comment says is written in Jacques's voice and must never go to her. This audience *is* her. **A partner-side sequence has to be written before a capture box belongs on that page.** That sequence does not exist yet and is the obvious next build.
- **The ACA Laundry List is not reproduced.** It's theirs; the page describes the traits in its own words and links to `adultchildren.org`. Same not-affiliated line as the Al-Anon page.
- `/alcoholic-personality` **opens by denying its own head term** — "there isn't one" — then explains what each behavior is *doing*. It is one keystroke from a page that labels people. Do not "strengthen" it.
(Two outbound links went out unverified: `al-anon.org` and `adultchildren.org` are both blocked from session environments. Worth one click each.)

■ LAST SEMRUSH PULL — the trial is ending, KEYWORDS.md is now closed

Jacques said the trial was nearly up, so a final sweep went into `KEYWORDS.md` ("FINAL PULL — 12 Aug 2026"). **After this there is no way to pull a number. Quote the file; never re-estimate, and never assume a term missing from it is unclaimed — it may simply never have been pulled.**

Three lanes came out of it that the file never had, all bigger than anything currently built:

1. "dry drunk" — ~11,700/mo at difficulty 30-34, one page. Sober and still the same person. It is the closest term in the whole dataset to what Jacques actually lived, the site never mentions it, and it is the strongest unbuilt page anywhere in the file.

2. Adult children of alcoholics — 8,100/mo at KD 31, plus the "laundry list" cluster at ~6,000/mo and difficulty 14-24. **A third audience,** not a variant of the drinker or the partner: the person who grew up in it.

3. "Alcoholic personality" — ~9,400/mo, all KD 33 or under, six phrasings of one question. Tone warning in the file: one keystroke from a page that labels people.

The Skip rule held again and it cost the most this time: `aca meetings` and `acoa meetings` near me are 14,000+/mo at difficulty 11-21 — the cheapest traffic in the document — and they are tagged Skip for the same reason as

al anon meetings near me. We cannot say where somebody's meeting is.

The finding that should steer marketing

Semrush, live: ****the partner audience is ~74,500 US searches a month against ~23,000 for the drinker.**** Three times bigger, and the segment nothing else in the category serves.

But the aim is wrong. The two words that audience actually types are al anon **(33,100/mo)** and ****codependency (33,100/mo, CPC \$0.12 — nobody monetizing it)****. Both now have pages, built later the same day — see the block above for what they actually target, which is not the head terms.

is-my-husband-an-alcoholic.html targets a **260/mo** phrase.

Best term already owned: hangxiety — 12,100/mo, difficulty 35, CPC \$0.11.

Three corrections to older numbers: "binge eating at night" is **140/mo**, not 3,480 (that was the whole cluster reported as one term). partner-watches-porn targets a **20/mo** phrase. **Gambling is difficulty 94 — dead as SEO**, still alive as video and outreach.

Outreach state

- **7 emails sent 12 Aug**, plus NCPG, Betrayal Trauma Recovery and SMART Recovery later the same day.
- **Missouri DBH replied** and named the door: their listings are limited to *"affiliated national organizations such as SAMHSA and NASADAD."* NASADAD was emailed the same morning, which makes it ****the highest-value pending item in the whole file.**** No reply owed to DBH now. The reply that *is* owed comes later, in that same thread, if NASADAD lists it: ****"You mentioned NASADAD. We're on their list now."**
- **Salvation Army St. Louis ARC — 3949 Forest Park Ave, (314) 535-0057.** Ten minutes from Jacques. Free residential rehab. This is a walk-in, not an email.
- **SHRM is two targets.** St. Louis chapter directory costs \$500/\$700 — skip. Their **speaker call is free** (7 meetings/yr, 200+ HR people). The national vendor directory is free but goes through a ****Calendly call with a MediaBrains rep****, not a form.

■ FIXED 12 Aug — Studio b0845, the Quick Video decimal bug

1.3s set a shot to **3 seconds**, 2.0s to **1**, and the row number turned up at the front of the caption. **That was one bug, not two:** the leading `/\d{1,3}\s*[. ]/` numbering strip ate the 1. of 1.3s as a row number and left 3s for the duration match to find.

The strip now requires the dot to be followed by a space or end of line — a row number is 1. , a decimal is 1.3 — and the duration regex accepts decimals to one place. Verified against ten lines through the real `parseShotLine()`, including the regressions that had to keep working (1. 30s A long hold. still reads 30; Day 400 was the hard one. still refuses to treat 400 as a length).

Still true, and not a bug: "■ Cut on the beat of the song" is on by default and rounds every shot up to a whole musical bar with a 1.5s floor, so sub-second times still come out longer with it ticked. Uncheck it for the episode slideshows. If it ever becomes worth fixing properly, the culprit is `Math.max(1, Math.round(d / bar)) * bar` — forcing a minimum of one full bar is what turns a 17-second piece into 30+ on a slow song.

Also: "■ Cut on the beat of the song" is **on by default** and snaps every shot to a whole musical bar with a 1.5s floor. On a sparse instrumental that turns a 17-second piece into 30+. Uncheck it for the episode slideshows.

CURRENT STATUS (2026-08-06, later) — read this before anything below

****Latest facts, verified from the code this session — these override anything older that disagrees:****

- **App version 5.1.0** (APP_VERSION, sw.js CACHE_NAME). ****Studio build b0807.**** Earlier notes saying 5.0.1 / b0804 / b0806 are stale.
- **Web push shipped (5.1.0):** lesson reminders now arrive with the app fully closed (server/push.js, new db tables). Also fixed: Reset Recovery keeping the old date, and lesson reminders never reaching Pro (5.0.3).
- **Phone number at signup is back** (commit 6bedd08, optional field, #gate-phone). It was present when Data safety was first completed, so the form still matches the app — ****Jacques confirms nothing on the Play side has been touched since the original submission, and nothing needs changing.**** Closed; don't re-raise it.
- **Studio, this session:** multi-select delete in My Media (■ Select multiple → tap thumbnails → bulk delete); **touch-drag timeline reorder** — the old reorder used the HTML5 `draggable` API, which mobile browsers never fire for a finger, so it only ever worked with a mouse; now Pointer Events on a ■ handle, verified with real touch input. Arrows kept as fallback.
- **Studio caption sizing fixed (b0807):** Studio has TWO text systems — server-burned ASS **captions** and client-drawn **overlays** (title cards and

Quick Video's per-picture words). The Size control only fed captions, and Quick Video's text was hardcoded at size 64 (~40% larger than a real caption, so even "Small" read big). Size now multiplies overlays too

(captionSizeMul), and auto-written text uses `AUTO_TEXT_SIZE = 0.042 * 1080`

— the same fraction as `ASS_SIZES.medium`. Confirmed working by Jacques.

- `reference/Getting-Noticed-Course.pdf` **(new)**: a 16-page sales/marketing course written for him — 8 modules, quizzes, final exam, reasoned answer key. Module 0 is how to work with Claude (repo-as-memory, "read START-HERE.md"). Meant for the Wednesday small-win slot, one module a week.

- **Also landed from other sessions:** `reference/SCHOOLS-BENEFITS-PLAYBOOK-2026-08.md` (the employer/school-benefits angle, with ready-to-send emails) and `reference/ADS-AND-CONTENT-MISTAKES-2026-08.md` (advertising and social content mistakes, with live paid-search data).

Note on rhythm: he moves between several concurrent sessions. Always

`git fetch origin main && git merge origin/main` before editing, and always re-read this block rather than trusting a summary in your own context.

"Does this change affect the Play Store test?" — the standing answer

He asks this on most changes. The app is a **TWA**: the `.aab` is an empty shell that opens `https://www.turnsomedayintodayone.com/app?src=play`. So:

Change	Play impact
Anything in <code>index.html</code> , <code>server/</code> , pages, lessons, audio, Studio	None. Ships with the Railway deploy. No <code>.aab</code> , no review, no clock reset.
A web API the page calls (wake lock, push, camera, notifications)	None for the .aab — the browser handles it, Android sees no new permission.
Anything in <code>TurnSomeDayIntoOneday/twa/</code> — package name, icon, splash, target SDK, <code>startUrl</code>	New .aab required. Bump <code>appVersionCode</code> , bubblewrap build, re-upload.
Anything that collects new user data	No .aab, but the Data safety form must be updated BEFORE it ships.

How to check, don't guess: `git log --oneline -- TurnSomeDayIntoOneday/twa/`

— if that path is untouched, no new build is needed. As of 2026-08-06 the shell is still at `appVersionCode: 1` and has not changed since it was first created.

Verified example: the SOS/lesson wake-lock fix (PR #46, branch

`claude/sos-talk-screen-cutoff-dfhjw4`) — web-only, TWA untouched, zero Play impact.

****2026-08-06** — Jacques tested the whole app on his real phone and reports everything working**, SOS "Talk me through it" and lesson audio included. He

has not touched the Play Store configuration since the very first submission, and the `twave/` folder confirms it (`appVersionCode` still 1, untouched since creation). **Nothing is outstanding on the Play side.** The only remaining blocker to production is the one that was always there: 12 testers opted in for 14 continuous days.

Phone number / Data safety — resolved 2026-08-06. The optional phone field was in the app when the Data safety form was originally completed; it was removed for a while and has since been restored, so the form still describes what the app actually does. ****Jacques confirms he has not changed anything on the Play side since the first submission, and nothing needs changing.**** Treat this as closed. (If a future session ever *adds* a genuinely new data type, that's when the form moves — not for a field that was always declared.)

STATUS (2026-08-06, earlier)

Everything under "Where things stand" is older background. This block is the live picture. **Update this block at the end of every session** — a session that reads only the older text will give him advice that's a week behind.

- **Play Store: still in CLOSED TESTING.** Not applied for production yet.

Testers still running the 14-day window. Don't touch billing; don't plan Play-side work; don't tell him to "check production."

- **Outreach: 10 of 30 sent** (Weeks 1 and 2 ticked in `OUTREACH.md`).

Week 3's five are next. He is doing the sales work — do not imply otherwise.

- **Content: live and posting.** 17 scripts loaded into Buffer; refills happening 6 Aug. Videos are out in the world. The bank is being worked through, not sitting idle.

- **Branch drift found 6 Aug — FIXED same day (his go-ahead):** the stranded `claude/app-qc-competitive-analysis-lehsn9` branch was fast-forwarded level with `main`, and the version quadruple (`sw.js` `CACHE_NAME`, `APP_VERSION`, both `package.json` files) was bumped **7.0.0 → 7.0.1** for the wake-lock change, per the `HANDOFF` rule. Keep both habits: every ship moves the quadruple and lands on both branches plus `claude/vibe-code-uwx1k` (deploy).

- **Studio: build stamp is now b0806** (the 08-05 notes still said b0804 — it has moved). Confirmed present in `Studio/web/index.html`: templates, the teleprompter, and webcam recording (`getUserMedia`). This is the toolchain `DO-THIS-NEXT.md` tells him to record with, so it is working end to end.

- **SOS + lessons screen-cutoff fix: merged to main and in this repo.** A shared

screen wake-lock manager (wakeLockAcquire/Release/Ensure, index.html ~6271) now serves the SOS voice guide, breathing, urge surfing, panic mode and lesson audio; re-acquires if the system revokes it; resumes on unlock. ****No Play impact**** — the Android app is a TWA pointing at the live site, so the fix ships with the server deploy: no new .aab, no version bump, no review, no reset of the 14-day clock, no new permissions. Only step left is the site being deployed from main. **Verified on a real phone by Jacques, 6 Aug — it works.**

Deploy path confirmed: Railway deploys `claude/vibe-code-uwxxlk`, which is at the same commit as `main` (92165e7), so the fix is live. Service-worker cache is not a blocker — `sw.js` serves pages network-first, so a new `index.html` lands on the next load even though the version quadruple wasn't bumped.

- *Caveat for a future session:* the write-up that shipped it said both features rely on the phone's built-in voice and that "true screen-off audio would need real audio files." That's not right — SOS defaults to real MP3s (`vgStartAudio`, phone TTS is only the fallback option) and all 395 lessons have MP3s. The audio itself can survive a lock; what stalls is the **setTimeout-driven step sequencing**, which browsers freeze on a hidden page. Keeping the screen on dodges that, so his symptom is fixed — but real screen-off playback is a solvable job (drive steps off `audio timeupdate/ended + MediaSession`), not an impossible one. Don't repeat the claim that it needs recordings he doesn't have.
- **6 Aug: KEYWORDS.md created at the repo root** (two verified Semrush pulls, final dataset — the marketer agent is unblocked), and the first two pages of its page plan are LIVE: **/partner-drinks** and **/partner-watches-porn** — the inclusive supporter pages (husband, wife, boyfriend, girlfriend named on each; KD 6-25 cluster). Routes, sitemap, landing-footer links, and a when-he-drinks cross-link all wired; crisis resources above every signup link; server boot-tested. Version 7.0.3. Also live 6 Aug: **/how-to-stop-binge-eating** (the 9,900/mo anchor), **/betrayal-trauma-recovery** (the moat's companion; cross-linked from **/partner-watches-porn**), and **/how-to-stop-drinking** (long-form; the cold-turkey SAFETY warning appears twice — above the fold AND above signup). **The KEYWORDS.md five-page plan is COMPLETE.** All five: routes, sitemap, landing footer, crisis-above-signup, server boot-tested. Version 5.0.1.
- **Version reset 7.0.3 → 5.0.0 on 6 Aug at Jacques' request.** Precedented (35→7, 12→7); the updater compares SHAs, the number only has to change.

Future bumps go 5.0.1, 5.0.2... — do not "correct" it back upward.

- **Open from 08-05, still unresolved:** the two origin-story picture fixes

(real app screenshot on the laptop shot; "Someday" split into two words), and the watermark "Patch it out" tool — it's ffmpeg de logo (pixel averaging), so it smears on busy backgrounds. Real fix would be AI inpainting; not built.

Where things stand (2026-08-03, historical)

Done & live: Pro chat cap set to 30/day everywhere; Pro "X of 30 left" counter; /when-he-drinks partner landing page; /go/tiktok|youtube|facebook tracked bio links (all three set in the socials); shared-device logout wipe; offline boot fix; guided-mode label removed; in-app updater fixed; SOS "Talk me through it" now real MP3 audio with a voice picker (survives screen lock); wake-lock keeps the screen awake during lessons; SEO groundwork shipped (sitemap.xml, robots.txt, homepage footer links, /best-recovery-apps roundup, 16 competitor "-alternative" pages). Google Search Console is verified and the sitemap submitted (done 2026-08-03)

- check it ~monthly for which searches show his pages; don't re-walk setup.

Lessons overhaul (2026-08-03, done in a parallel session, merged to main):

the app now targets EVERYONE - straight, gay, bi, gender-neutral, men and women; all lessons rewritten inclusive/gender-neutral and claims-audited. Recorded audio for all 395 lessons x 5 voices (1,975 MP3s) with a voice picker like SOS; pause button (no more restart-from-zero); pop-out big player with read-along; on completion auto-scroll + next lesson (follows multi-addiction tracks); audio keeps playing with the screen off. Find-anything button moved off the Friendly send button; phone-number option removed.

Division of labor (2026-08): Jacques' wife runs social comments and followers

- he never opens the feeds or dashboards; his daily is 10 min (Buffer check).
- Positioning widened by his call: post for EVERYBODY (HIM, HER, alcohol, food), Monday stats decide the lean - nothing says "her first" anymore.

Reality check (as of 08-03): /admin/stats shows only Jacques' own test signups - normal; first outreach batch and consistent Buffer posting just starting. First stranger signup expected in weeks, not days. Don't let a session panic about flat numbers. Cleanup pending: remove test123 + test signups when he gives the word (keep his real owner account).

Studio (2026-08): Dynamics templates shipped - 10 free CapCut-style one-tap effects (Beat Mix, B/W Strobe, Color Flip, Thermal, Glow Up, Echo Trails, Hype Shake, Mirror World, Time Machine old-to-new, 3-2-1 Opener) on any video OR

picture (stills become 4-15s clips). Build stamp b0804.

Deferred (post-launch, see build list):

1. ~~Lesson background audio~~ DONE 2026-08-03 (all 395 lessons, 5 voices, free).
2. Wire up real Friendly AI (ANTHROPIC_API_KEY on the server) — also activates the Pro chat counter (guided replies don't consume quota, so it reads full now).
3. Accountability partner link (/together) — needs Data safety update.

Marketing next actions: make videos in Studio → post via Buffer 3×/day (Morning=HIM, Midday=HER, Night=PROOF) → reply to comments → check /admin/stats Mondays → do more of what grows. (Superseded 08-04: post for EVERYBODY, the Monday numbers pick the lane — nothing leads with "her" anymore.)

Standing reminder: a weekly Routine fires into the owning session every Wednesday to handle the Play-side update + offer the deferred features the moment the app is live in the Play Store.

Pictures and video — who can fetch what

Recorded 12 Aug 2026 because a session tried to walk Jacques through setting up something he had already been using for weeks.

- **Pexels is already set up in Studio, and has been.** Studio → My Media →
■ **Free stock b-roll**, with ■ **Clips** and ■ **Photos** chips. Free key, free clips, no attribution, imports straight to My Media and onto the timeline. **Do not tell him to go and get a key — he has one.** The key lives in his local .env and is deliberately not in this repo.
- **This is the first place to look for anything that isn't a scripted moment** — establishing shots, weather, textures, transitions. An AI scene is ~\$2; a Pexels clip is free and lands in two taps. Reserve fal.ai spend for shots Pexels cannot give.
- **A session CANNOT search Pexels.** pexels.com and api.pexels.com are both blocked by the session network policy, there is no Pexels connector in the registry, and the key isn't here. Verified 12 Aug. A session's job is to write the search terms; Jacques runs the search in Studio.
- **A session CAN search Shutterstock** (connector, read-only: previews and metadata, no licensing or download). It's paid, and it has no video import path into Studio — so it's the fallback for one specific shot Pexels lacks, not the default.
- **Faces are the thing to avoid, on either service.** Shutterstock's license restricts showing a recognizable person in connection with sensitive subjects

(addiction, mental health) without a "posed by a model" disclaimer, and Pexels' license says identifiable people may not appear "in a bad light or in a way that is offensive" — with no disclaimer option at all. This app's whole subject sits inside that clause. **Search rooms, weather, hands, doors, roads, light — not people.** The text cards already work precisely because the reader supplies the face.

- **What a session can and cannot make.** No session can generate an image or a video — there is no image model here. What it CAN do is write code that draws deterministically (reference/make-song-cards.py is PIL: gradient, icon, type), and write the prompts that Studio sends to fal.ai. Cards are free and instant; AI scenes cost money and run in Studio, on his account.
 - Licence terms for everything in use are logged in reference/asset-licenses-2026-08-08.md.
-

Rules of the house

- **Every email drafted into Gmail is written as HTML with a real link, never as a bare URL in plain text.** Gmail rewrites every link it sends through google.com/url?q=...&source=gmail&ust=... — that happens to everyone, on every message, whether the link was typed by hand or inserted through the API, and it cannot be turned off. What it *can* be stopped from doing is making that string the text the reader sees. Write `<a href="https://www.turnsomedayintodayone.com">www.turnsomedayintodayone.com</a>` and the recipient reads the address; the redirect stays in the invisible href where it belongs. Write the bare URL and the reader gets a wall of google.com/url?q= in the middle of a cold pitch, which reads like tracking. Learned 11 Aug 2026, after 22 school emails and two podcast pitches had already gone out the wrong way. **Broken again 12 Aug** — NCPG, Betrayal Trauma Recovery and SMART Recovery were all drafted with a bare URL and sent before it was caught. Twice now. **Use `htmlBody` with a real `<a href>` on every single draft**, including the ones that feel too short to bother with.
- **Jacques's send address is** turnsomedayintodayone@gmail.com. Not the address on the account profile. Every draft on 12 Aug signed off with the wrong one and had to be corrected — a reply landing in a different inbox than the one it came from is a reply that gets lost.
- **Put answers in chat, not only in a file.** Jacques does not read the repo and git is not installed on his machine. Writing something to a markdown file

and telling him the filename means he never sees it. Write it in the message, *then* also save it. He said so plainly on 12 Aug after a full set of episode captions went into a file he never opened.

- **Don't decide what to skip.** Parking a target, dropping a lane, or calling something not worth doing is his call, not the session's. On 12 Aug the Salvation Army was filed as "worth doing, not worth doing badly — park it," and it turned out to be the strongest target in the file and ten minutes from his house. If something is hard, do it or list it as undone with what is blocking it. No verdicts.
- Story, not statistics (see medical-claims-audit.md).
- **Music: offer a range, never pick one lane for him.** Corrected by Jacques himself on 12 Aug: "i do rnb and hiphop but i like a variety even latin music."* The old wording here said "never default to R&B, hip-hop or soul," which read as though he didn't want them — wrong. **He does R&B and hip-hop.** The actual problem is that tools keep *defaulting* to that lane for him without asking. So: offer a spread — R&B, hip-hop, **Latin**, folk, rock, country, cinematic, electronic — and let him choose. Same for any other assumption about his taste. Offer the range, don't choose for him. (Worth knowing: Episode 2's ROSA is Latina, so a Latin bed there is the right music for whose story it is, not decoration.)
- Privacy is the product: never expose one user's data to another; "he doesn't see what you write" must stay true.
- Anything worth remembering across sessions goes in a file — chat doesn't persist.